

**Московский государственный технический  
университет им. Н.Э. Баумана**

Факультет «Информатика и системы управления»  
Кафедра ИУ5 «Системы обработки информации и управления»

Курс «Теория машинного обучения»

Отчет по лабораторной работе №5  
«Ансамбли моделей машинного обучения»

Выполнил:  
студент группы ИУ5-63Б  
Ветошкин Артём

Подпись и дата:  
14.06.22

Проверил:  
Юрий Евгеньевич Гапанюк

Подпись и дата:

Москва, 2022 г.

# Ансамбли моделей машинного обучения.

## Импорт библиотек

```
import numpy as np
import pandas as pd
from io import StringIO
from IPython.display import Image
from heamy.dataset import Dataset
from heamy.estimator import Regressor, Classifier
from heamy.pipeline import ModelsPipeline
from gmdhpy import gmdh
import pydotplus
import seaborn as sns
import xgboost as xgb
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import make_scorer
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeRegressor, export_graphviz
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from sklearn.svm import SVR
import operator
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.datasets import *
sns.set(style="ticks")
```

## Загрузка данных

В качестве набора данных мы будем использовать набор данных о заболеваемости диабетом из пакета sklearn. Загрузим файлы датасета из библиотеке `sklearn.datasets`, и преобразуем в `DataFrame` библиотеки `Pandas`.

```
ds = load_diabetes()
df = pd.DataFrame(data= np.c_[ds['data'], ds['target']],
                  columns= list(ds['feature_names']) + ['target'])
df.count()
```

```
age      442
sex      442
bmi      442
bp       442
s1       442
s2       442
s3       442
s4       442
s5       442
s6       442
target   442
dtype: int64
```

```
df.corr()
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```

	age	sex	bmi	bp	s1	s2	s3	s4	s5	s6	target
age	1.000000	0.173737	0.185085	0.335427	0.260061	0.219243	-0.075181	0.203841	0.270777	0.301731	0.187889
sex	0.173737	1.000000	0.088161	0.241013	0.035277	0.142637	-0.379090	0.332115	0.149918	0.208133	0.043062
bmi	0.185085	0.088161	1.000000	0.395415	0.249777	0.261170	-0.366811	0.413807	0.446159	0.388680	0.586450
bp	0.335427	0.241013	0.395415	1.000000	0.242470	0.185558	-0.178761	0.257653	0.393478	0.390429	0.441484
s1	0.260061	0.035277	0.249777	0.242470	1.000000	0.896663	0.051519	0.542207	0.515501	0.325717	0.212022
s2	0.219243	0.142637	0.261170	0.185558	0.896663	1.000000	-0.196455	0.659817	0.318353	0.290600	0.174054
s3	-0.075181	-0.379090	-0.366811	-0.178761	0.051519	-0.196455	1.000000	-0.738493	-0.398577	-0.273697	-0.394789
s4	0.203841	0.332115	0.413807	0.257653	0.542207	0.659817	-0.738493	1.000000	0.617857	0.417212	0.430453
s5	0.270777	0.149918	0.446159	0.393478	0.515501	0.318353	-0.398577	0.617857	1.000000	0.464670	0.565883
s6	0.301731	0.208133	0.388680	0.390429	0.325717	0.290600	-0.273697	0.417212	0.464670	1.000000	0.382483
target	0.187889	0.043062	0.586450	0.441484	0.212022	0.174054	-0.394789	0.430453	0.565883	0.382483	1.000000

```
df.isnull().sum()
```

```
age      0
sex      0
bmi      0
bp       0
s1       0
s2       0
s3       0
s4       0
s5       0
s6       0
target   0
dtype: int64
```

## Подготовка данных

Поделит данные на тестовую и тренировочную выборку в пропорции 1 к 4

```
df_X_train, df_X_test, df_y_train, df_y_test = train_test_split(
    df.drop(columns='target'), df['target'], test_size=0.2, random_state=12)
```

## Случайный лес

```
tree = RandomForestRegressor(n_estimators=10, random_state=12)
tree.fit(df_X_train, df_y_train)
tree_y = tree.predict(df_X_test)
```

```
tree_y
```

```
array([184.2, 136.3, 143.7, 114.2, 194.2, 160. ,  71.1, 245.4, 136.7,
       243.8, 260.5, 201.4, 161. , 154.4, 105.5, 224.8, 173.3,  66.5,
       199. , 152.2, 156.2, 122. , 261.1, 142.5, 157.8, 108.8, 156. ,
```

```
126.6, 68. , 144.4, 174.4, 162.7, 181.5, 93.4, 117.1, 191.8,
112.1, 212.8, 222.9, 191. , 236.7, 200.5, 123.4, 77. , 182.3,
76.6, 223.3, 80.9, 263.9, 290.8, 95.8, 172.5, 214.2, 181.3,
150.5, 132.3, 142.8, 205.9, 225. , 238.1, 266.8, 101. , 140.8,
113.8, 76.5, 71.1, 139.7, 216.7, 100.7, 155.6, 85.8, 72.2,
79.9, 82. , 123.8, 190.6, 126.7, 129.6, 152.8, 181.6, 117.5,
179. , 91.2, 84.9, 100.2, 174.3, 244.2, 148.2, 56.2])
```

## Бустинг

```
reg = xgb.XGBRegressor(
    n_estimators=10,
    eval_metric=mean_squared_error,
)
reg.fit(df_X_train, df_y_train, eval_set=[(df_X_train, df_y_train)])
reg_y = reg.predict(df_X_test)
```

```
[0] validation_0-rmse:126.78217
[1] validation_0-rmse:95.65629
[2] validation_0-rmse:73.21293
[3] validation_0-rmse:57.39291
[4] validation_0-rmse:46.06264
[5] validation_0-rmse:38.16575
[6] validation_0-rmse:32.10063
[7] validation_0-rmse:27.78214
[8] validation_0-rmse:23.98275
[9] validation_0-rmse:21.76612
```

D:\anaconda3\lib\site-packages\xgboost\data.py:250: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas in a future version. Use pandas.Index with the appropriate dtype instead.

```
elif isinstance(data.columns, (pd.Int64Index, pd.RangeIndex)):
```

```
reg_y
```

```
array([149.30428 , 105.880424, 145.81374 , 102.04156 , 174.46553 ,
       114.801186, 110.59888 , 243.877 , 154.5212 , 237.16719 ,
       222.18236 , 217.16653 , 150.85403 , 134.93927 , 95.98342 ,
       237.88918 , 180.51299 , 83.54547 , 207.24379 , 109.5174 ,
       142.41653 , 119.48074 , 239.61699 , 103.94377 , 155.80708 ,
       130.35544 , 102.680244, 89.2868 , 76.83893 , 145.3199 ,
       104.73179 , 145.29933 , 134.9832 , 120.21382 , 85.30003 ,
       168.91045 , 76.302284, 209.41447 , 228.50658 , 179.25017 ,
       202.60506 , 186.886 , 85.2626 , 101.48372 , 173.86761 ,
       73.36852 , 214.41394 , 84.811356, 232.04376 , 272.86478 ,
       101.107605, 146.30386 , 185.52725 , 171.97969 , 142.88792 ,
       112.13408 , 123.32981 , 248.35106 , 223.41351 , 244.65875 ,
       247.1172 , 69.80298 , 153.20953 , 87.67907 , 92.2524 ,
       83.07777 , 85.56129 , 217.9154 , 94.80268 , 129.47073 ,
       83.429146, 83.567856, 82.01156 , 52.425613, 99.1951 ,
       143.71127 , 154.04832 , 115.452675, 132.10777 , 186.45786 ,
       102.96815 , 208.654 , 74.69237 , 78.96792 , 96.96595 ,
       137.61682 , 213.70499 , 121.17906 , 57.572853], dtype=float32)
```

## Стекинг

```
dataset = Dataset(df_X_train, df_y_train, df_X_test)

model_rf = Regressor(dataset=dataset, estimator=RandomForestRegressor, parameters={'n_estimators': 10}, name='rf')
model_lr = Regressor(dataset=dataset, estimator=LinearRegression, parameters={'normalize': True}, name='lr')
```

```
pipeline = ModelsPipeline(model_rf, model_lr)
stack_ds = pipeline.stack(k=15, seed=111)

stacker = Regressor(dataset=stack_ds, estimator=DecisionTreeRegressor)
stacker_y = stacker.predict()
results = stacker.validate(k=15, scorer=mean_squared_error)
```

Metric: mean\_squared\_error  
Folds accuracy: [7022.25, 12363.958333333334, 5662.666666666667, 6137.416666666667, 8583.041666666666, 7334.791666666667, 5365.958333333333, 7370.333333333333, 6070.95652173913, 4447.173913043478, 4680.826086956522, 4341.521739130435, 8708.91304347826, 4379.521739130435, 4062.6521739130435]  
Mean accuracy: 6435.465458937198  
Standard Deviation: 2163.197180008688  
Variance: 4679422.03959754

stacker\_y

array([237., 170., 302., 137., 189., 124., 79., 129., 244., 212., 277.,  
288., 263., 103., 78., 275., 214., 59., 85., 144., 201., 55.,  
243., 144., 60., 50., 262., 200., 128., 71., 91., 109., 138.,  
68., 37., 222., 137., 233., 258., 172., 257., 192., 150., 114.,  
118., 129., 197., 64., 258., 245., 60., 68., 189., 235., 185.,  
172., 198., 233., 275., 164., 242., 131., 110., 87., 45., 63.,  
136., 232., 160., 272., 137., 64., 90., 65., 78., 131., 277.,  
113., 164., 265., 71., 110., 135., 91., 200., 128., 173., 78.,  
104.]])

Многослойного персептрона

Для реализации данной модели будем использовать функционал библиотеки TensorFlow

```
# Нормализуем обучающую выборку с помощью отдельного слоя нормализации
normalizer = tf.keras.layers.Normalization(axis=-1)
normalizer.adapt(np.array(df_X_train))
```

print(normalizer.mean.numpy())

[[ -0.00158359 -0.00089624 0.00031588 -0.00096308 -0.00051653 -0.00046802  
 -0.00057182 0.00042046 0.0001608 0.00014264]]

df\_X\_train.describe()

```
.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}
```

	age	sex	bmi	bp	s1	s2	s3	s4	s5	s6
count	353.000000	353.000000	353.000000	353.000000	353.000000	353.000000	353.000000	353.000000	353.000000	353.000000

	age	sex	bmi	bp	s1	s2	s3	s4	s5	s6
mean	-0.001584	-0.000896	0.000316	-0.000963	-0.000517	-0.000468	-0.000572	0.000420	0.000161	0.000143
std	0.049006	0.047567	0.049150	0.048559	0.049247	0.048940	0.048204	0.048292	0.047684	0.046903
min	-0.107226	-0.044642	-0.090275	-0.112400	-0.126781	-0.112795	-0.102307	-0.076395	-0.126097	-0.129483
25%	-0.041840	-0.044642	-0.036385	-0.036656	-0.035968	-0.032002	-0.036038	-0.039493	-0.033249	-0.030072
50%	0.005383	-0.044642	-0.008362	-0.005671	-0.004321	-0.005072	-0.006584	-0.002592	-0.003304	-0.001078
75%	0.038076	0.050680	0.033673	0.034496	0.027326	0.031254	0.026550	0.034309	0.033047	0.027917
max	0.110727	0.050680	0.170555	0.132044	0.153914	0.198788	0.181179	0.185234	0.133599	0.135612

```
#Создадим нейронную сеть с одним скрытым слоем
nn_model = keras.Sequential([
    normalizer,
    layers.Dense(20, activation='relu'),
    layers.Dense(1)
])

nn_model.compile(loss='mean_squared_error',
                  optimizer=tf.keras.optimizers.Adam(0.001))
history = nn_model.fit(
    df_X_train,
    df_y_train,
    validation_split=0.2,
    verbose=0,
    epochs=400
)
nn_y = nn_model.predict(df_X_test)
nn_y
```


3/3 [=====] - 0s 1ms/step

```
array([[212.92929 ],
       [ 19.594112],
       [106.97292 ],
       [ 70.99605 ],
       [189.48883 ],
       [ 86.36806 ],
       [111.75647 ],
       [116.60968 ],
       [122.602585],
       [186.56653 ],
       [243.88792 ],
       [152.87848 ],
       [127.72024 ],
       [136.28882 ],
       [ 60.78453 ],
       [230.90665 ],
       [129.89294 ],
       [106.85017 ],
       [167.22697 ],
       [150.80173 ],
       [ 68.531204],
       [142.75902 ],
       [283.63235 ],
       [128.60912 ],
       [152.00255 ],
       [118.70987 ],
       [114.60079 ],
       [ 63.414852],
       [ 66.22781 ],
       [ 60.62016 ],
```

```
[163.18936 ],
[125.61765 ],
[152.70609 ],
[ 66.165245],
[ 51.140064],
[192.25508 ],
[ 86.394844],
[181.1155  ],
[172.6566  ],
[110.59793 ],
[281.3926  ],
[225.30415 ],
[ 83.13081 ],
[ 92.2989  ],
[199.2587  ],
[ 74.624916],
[160.40106 ],
[ 58.419716],
[235.88925 ],
[310.13818 ],
[101.25036 ],
[111.54132 ],
[198.2014  ],
[121.38501 ],
[172.93774 ],
[ 91.53825 ],
[ 71.94856 ],
[214.13223 ],
[174.2966  ],
[196.82266 ],
[273.973  ],
[ 65.883736],
[134.82565 ],
[ 92.42822 ],
[153.13995 ],
[101.361725],
[164.9628  ],
[202.02533 ],
[ 87.275406],
[135.85161 ],
[ 31.933609],
[ 85.35344 ],
[105.59263 ],
[ 27.264675],
[120.45435 ],
[120.036194],
[146.60416 ],
[ 42.239822],
[128.82458 ],
[177.50784 ],
[138.09705 ],
[108.38889 ],
[ 49.237446],
[ 43.287865],
[122.67038 ],
[186.49951 ],
[278.9737  ],
[ 97.3726  ],
[109.22105 ]], dtype=float32)
```

```
# посмотрим историю обучения
```

```
plt.plot(history.history['loss'], label='loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.xlabel('Epoch')
plt.ylabel('Error')
plt.legend()
plt.grid(True)
```



## Модель МГУА

```
gmdh_model = gmdh.Regressor(ref_functions=('linear_cov', 'quadratic', 'cubic', 'linear'))
gmdh_model.fit(df_X_train.values, df_y_train.values)
gmdh_y = gmdh_model.predict(df_X_test.values)
gmdh_y
```

```
train layer0 in 0.18 sec
train layer1 in 0.53 sec
train layer2 in 0.52 sec
train layer3 in 0.52 sec
train layer4 in 0.51 sec
train layer5 in 0.53 sec
```

```
array([209.10020708, 122.5295096 , 159.51721852, 99.38626424,
       193.47445709, 138.60592576, 122.73902489, 165.22235386,
       178.97968777, 189.02133524, 258.40734395, 179.30405188,
       135.75473228, 141.39286193, 68.37846824, 218.25632299,
       152.71693688, 73.28522292, 171.01202591, 194.23892706,
       127.43402267, 81.2616121 , 244.41066279, 142.97809548,
       148.28352379, 105.93573428, 164.23881249, 116.73649724,
       91.84712255, 127.59577046, 153.39128262, 178.53827268,
       176.22114935, 126.20823383, 103.85699124, 169.48493727,
       101.92350582, 194.07724585, 168.23861498, 145.92825908,
       264.6936559 , 230.78248817, 142.74540707, 84.90167377,
       187.76470783, 124.70535705, 179.84582544, 103.29003094,
       199.52872972, 271.75801822, 97.38062031, 124.40470526,
       211.20313612, 155.36240341, 197.9754771 , 156.52174031,
       173.60673959, 187.08229415, 189.81847303, 194.86499065,
       249.71941818, 140.51922649, 128.79353257, 94.69480909,
       66.84532263, 58.24483861, 165.47981933, 204.76503817,
       109.4112695 , 164.12303524, 124.14778806, 107.53705998,
       129.66411255, 90.37030001, 139.87412475, 170.05686851,
       188.03750376, 133.26186746, 170.62774333, 211.73070802,
       143.08267762, 163.43940807, 103.99513899, 100.92269545,
       140.97686736, 166.28475172, 248.4229277 , 157.79975723,
       65.79076495])
```

## Оценки моделей

```
results_metrics = [mean_squared_error(df_y_test, tree_y), mean_squared_error(df_y_test, reg_y),
mean_squared_error(df_y_test, stacker_y), mean_squared_error(df_y_test, nn_y), mean_squared_error(df_y_test, gmdh_y)]
model_list = ['random_forest', 'boosting', 'stacker', 'nn', 'gmdh']

sorted_el = list(sorted(list(zip(model_list, results_metrics)), key=lambda x: -x[1]))
results_metrics = list(map(lambda x: x[1], sorted_el))
model_list = list(map(lambda x: x[0], sorted_el))
```



```
fig, ax = plt.subplots(figsize=(20,20))
pos = np.arange(len(model_list))
rects = ax.barh(pos, results_metrics,
                 align='center',
                 height=0.5,
                 tick_label=model_list)
ax.set_title('mean_squared')
for a, b in zip(pos, results_metrics):
    plt.text(max(results_metrics)/2, a-0.05, str(round(b,6)), color='black')
plt.show()
```

